

NUR Hand-in Exercise 1

Angèl Pranger

February 26, 2025

Abstract

This document shows my solutions to hand-in exercise 1 of numerical methods in astrophysics.

1 Poisson distribution

For this exercise, I was asked to write a function that returns the Poisson probability distribution $P_\lambda(k) = \frac{\lambda^k e^{-\lambda}}{k!}$ for integers k given a positive mean λ .

It seems there are three potential problems to tackle. Firstly, there is the possibility of underflow for large values of k , caused by the division by a very large number $k!$. Secondly, there is the factor λ^k , which can cause overflow for large values of λ and/or k . Thirdly, there is the factor $e^{-\lambda}$, which can become too small for large values of λ .

I implemented three functions, which all compute the Poisson probability distribution in a different way. The first function only uses a conversion to log space, such that

$$\begin{aligned} P_\lambda(k) &= \exp\left(\ln\left(\frac{\lambda^k e^{-\lambda}}{k!}\right)\right) \\ &= \exp(\ln(\lambda^k) + \ln(e^{-\lambda}) - \ln(k!)) \\ &= \exp(k \ln(\lambda) - \lambda - \sum_{i=0}^{k-1} \ln(k-i)). \end{aligned}$$

The second function first uses not log space but a different order of operations, given by

$$P_\lambda(k) = e^{-\lambda} \cdot \frac{\lambda}{k} \cdot \frac{\lambda}{k-1} \cdot \dots \cdot \frac{\lambda}{1}.$$

The third function uses a combination of this different order of operations and log space, such that

$$P_\lambda(k) = \exp\left(\ln\left(\frac{\lambda}{k} \cdot \frac{\lambda}{k-1} \cdot \dots \cdot \frac{\lambda}{1}\right) - \lambda\right).$$

The script in which these functions are implemented is:

```
1 import numpy as np
2
3 # Using log conversion
4 def poisson1(lamda:np.float32, k:np.float32) -> np.float32:
5     """Computing the poisson distribution value for parameters (lamda, k) using log
6     conversion."""
7     j = np.float32(0)
8     for i in range(np.int32(k)):
9         i = np.float32(i)
10        j += np.float32(np.log(k-i))
11    P = np.float32(np.exp(k*np.log(lamda)-lamda-j))
12    return P
13
14 # Using different order of operations
15 def poisson2(lamda:np.float32, k:np.float32) -> np.float32:
16     """Computing the poisson distribution value for parameters (lamda, k) using an
17     unconventional order of operations."""
```

```

16 P = np.float32(np.exp(-lamda))
17 for i in range(np.int32(k)):
18     i = np.float32(i)
19     P *= np.float32(lamda/(k-i))
20 return P
21
22 # Using different order of operations and log conversion, avoiding np.exp(-lamda)
23 def poisson3(lamda:np.float32, k:np.float32) -> np.float32:
24     """Computing the poisson distribution value for parameters (lamda, k) using a
25     different order of operations and a log conversion."""
26     P = np.float32(1)
27     for i in range(np.int32(k)):
28         i = np.float32(i)
29         P *= np.float32(lamda/(k-i))
30     P = np.float32(np.exp(np.log(P) - lamda))
31     return P
32
33 # Computing and printing results
34 lamda_k = np.array([[1,0],[5,10],[3,21],[2.6,40],[100,5],[101,200]], dtype=np.float32)
35 print(" [lamda k] P1 P2 P3")
36 for values in lamda_k:
37     P1 = poisson1(values[0], values[1])
38     P2 = poisson2(values[0], values[1])
39     P3 = poisson3(values[0], values[1])
40     print(f" [{values[0]:.1f} {values[1]:.0f}] {P1:.6f} {P2:.6f} {P3:.6f}")

```

1.poisson.distribution.py

The results are given by the following. The first column states the values for λ and k , the columns named $P1, P2, P3$ give the results for the first, second and third function respectively.

```

1 [lamda k] P1 P2 P3
2 [1.0 0] 0.367879 0.367879 0.367879
3 [5.0 10] 0.0181328 0.0181328 0.0181328
4 [3.0 21] 1.01934e-11 1.01934e-11 1.01934e-11
5 [2.6 40] 3.61508e-33 3.61512e-33 3.61513e-33
6 [100.0 5] 3.10006e-36 3.15292e-36 3.10006e-36
7 [101.0 200] 1.26953e-18 1.07434e-05 1.26953e-18

```

poisson.distribution.txt

We see that the results of the three functions are the same up to 6 significant digits for the first three values of λ, k . This is to be expected, as the corresponding values of λ and k are relatively small and thus no over or underflow is expected for any of the functions. For the last two rows, where λ is relatively large, we find that the first and third function give the same result, while the result of the second function deviates. As potential under and overflow errors should not be exactly the same for the first and third function, this leads us to think that these functions actually give the correct answer (at least up to these 6 significant digits), while the second function then gives an incorrect result. The problem for the second function is here in the factor $e^{-\lambda}$, which for $\lambda = 100$ is of order 10^{-44} and can thus not be stored in a 32 bit float. The values $\lambda, k = 2.6, 40$ gives a slightly different result for each of the three functions. According to WolframAlpha, the result of the second function is actually the most accurate here, after which the result of the third function is a close second. It is not clear to us why the result of the first function deviates here, but this might have something to do with the fact that λ is relatively small while k is relatively large (we found in class that this value of k already gives problems with the original implementation of the Poisson distribution), or with the fact that $\lambda = 2.6$ is not an integer.

In conclusion, the third function gives overall the best result.

2 Vandermonde matrix

The code of any shared modules for question 2 is given by:

```

1 import numpy as np
2 import timeit
3 import sys

```

```

4 import os
5 import matplotlib.pyplot as plt
6
7 # Importing data
8 data=np.genfromtxt(os.path.join(sys.path[0],"Vandermonde.txt"),comments='#',dtype=np.
    float64)
9 x=data[:,0]
10 y=data[:,1]
11 xx=np.linspace(x[0],x[-1],1001) # x values to interpolate at
12
13 # Functions used
14 def LU_decomposition(matrix:np.ndarray):
15     """Takes a matrix and computes its LU-decomposition.
16     Returns the same matrix which now holds the LU-decomposition: the upper triangle and
17     diagonal hold the beta-ij (U),
18     while the lower triangle hold the alpha-ij (L). The alpha-ii are chosen to be equal
19     to 1."""
20     # Choose alpha-ii = 1
21     # Loop over the columns j
22     for j in range(matrix.shape[0]):
23         # Keep beta_0j = a_0j
24         # Update beta-ij where 0<i<=j
25         for i in range(1, j+1):
26             term = np.float64(0)
27             for k in range(0, i):
28                 term += matrix[i,k]*matrix[k,j]
29             matrix[i,j] -= term
30         # Update alpha-ij where i>j
31         beta_jj = 1/matrix[j,j]
32         for i in range(j+1, matrix.shape[0]):
33             term = np.float64(0)
34             for k in range(0, j):
35                 term += matrix[i,k]*matrix[k,j]
36             matrix[i,j] = (matrix[i,j]-term)*beta_jj
37     return np.float64(matrix)
38
39 def solving_system_with_LU(LU:np.ndarray, b:np.array) -> np.array:
40     """Solves the system Ax=b by taking the LU decomposition of A and b as input and
41     applying forward and backward substitution.
42     Returns the vector b which now holds the solution x."""
43     # Knowing that alpha-ii = 1
44     # Forward substitution
45     # Keep y_0 = b_0 as alpha_00 = 1
46     for i in range(1, LU.shape[0]):
47         term = np.float64(0)
48         for j in range(i):
49             term += LU[i,j]*b[j]
50         b[i] -= term
51     # Backward substitution
52     b[-1] /= LU[-1,-1]
53     for i in range(LU.shape[0]-1):
54         i = LU.shape[0]-1 - i
55         term = np.float64(0)
56         for j in range(i+1, LU.shape[0]):
57             term += LU[i,j]*b[j]
58         b[i] = (b[i]-term)/LU[i,i]
59     return np.float64(b)
60
61 def bisection(M:np.int64, x_samp:np.array, x_interp:np.array) -> np.ndarray:
62     """Takes as input a number of points M, a sample array x and an interpolation values
63     array x_interp.
64     Uses bisection to find the closest one or two points in x_samp for each point in
65     x_interp. Subsequently finds the closest M
66     points in x_samp and returns the index in x_samp of the lowest (leftmost) point."""
67     # Assumes x_samp are uniformly spaced
68     if (M > x_samp.shape[0]):
69         print("M is too large for the array x_samp")
70         return
71     # Loop over all points x in x_interp
72     for (i,x) in enumerate(x_interp):

```

```

68     # Set left and right edge indices within x_samp
69     l_idx = np.int64(0)
70     r_idx = np.int64(x_samp.shape[0]-1)
71     # Half the searching area, check in which half x lies, update edge indices for
    this half,
72     # until the left and right edge index differ by 0 or 1
73     while (r_idx - l_idx > 1):
74         m_idx = np.int64(np.floor((l_idx + r_idx)*0.5))
75         if (x <= x_samp[m_idx]):
76             r_idx = m_idx
77         elif (x > x_samp[m_idx]):
78             l_idx = m_idx
79     # Save the middle value of the current interval to later check whether x is
    closer to the left or right edge
80     m = (x_samp[l_idx] + x_samp[r_idx])*0.5
81     # Computing the left boundary index (ignoring range)
82     l_idx -= np.int64(np.ceil(0.5*M-1))
83     # Correcting left boundary if x is closer to the right boundary
84     # (in which case the right boundary is seen as the first point instead of the
    second)
85     if ((x > m) & (M % 2 == 1)):
86         l_idx += np.float64(1)
87     # Checking range (ensuring the left and right edge do not fall outside x_samp)
88     if l_idx < 0:
89         l_idx = 0
90     if (l_idx + M >= x_samp.shape[0]):
91         l_idx -= l_idx + M - x_samp.shape[0]
92     x_interp[i] = l_idx
93     return np.int64(x_interp)
94
95 def neville(M:np.int64, x_samp:np.array, y_samp:np.array, x_interp:np.array) -> np.
    ndarray:
96     """Takes an integer number of points M, sample array x_samp with corresponding y
    values y_samp and interpolation values x_interp.
97     Computes for each x in x_interp the value of the unique Lagrange polynomial through
    the M points in x_samp closest to x, using Neville's algorithm.
98     Returns the M-1 order Lagrange polynomial values corresponding to the x in x_interp
    and the corresponding error estimates."""
99     if (M < 1):
100         print("M is too small")
101         return
102     # Finding left index of M tabulated points using bisection
103     l_indices = bisection(M, x_samp.copy(), x_interp.copy())
104     # Initializing matrix P for Neville's algorithm (each row corresponds to a x in
    x_interp, each column corresponds to one of the M closest sample points)
105     Ps = np.zeros((x_interp.shape[0], M), dtype=np.float64)
106     for i in range(M):
107         Ps[:, i] = y_samp[l_indices+i]
108     # Looping over orders k and current intervals
109     for k in range(1, M):
110         for i in range(M-k):
111             j = i+k # note this j is only for the x_samp, as Ps is overwritten so just j
    =i+1 as index
112             # Ps[:, i] are overwritten to hold the linear interpolation between the
    previous values
113             Ps[:, i] = ((x_samp[l_indices+j]-x_interp)*Ps[:, i]+(x_interp-x_samp[l_indices
    +i])*Ps[:, i+1])/(x_samp[l_indices+j]-x_samp[l_indices+i])
114             if (M < 2): # No error estimate
115                 return Ps[:, 0], 0
116             error_estimate = np.absolute(Ps[:, 0] - Ps[:, 1])
117             return Ps[:, 0], error_estimate
118
119 def polynomial(c:np.array, x:np.array) -> np.array:
120     """Constructing the values y of the polynomial with coefficients given by c at
    points x.
121     Returns y values."""
122     y = np.float64(0)
123     for j in range(c.shape[0]):
124         y += c[j] * np.power(x, j)
125     return np.float64(y)

```

2.1 a

For this exercise, I wrote a code that firstly constructs the Vandermonde matrix V for the 20 data points x using that $V_{ij} = x_i^j$. Subsequently, the LU decomposition of V was found and the system $Vc = y$, for which the data points y were given, was solved for c . The coefficients c of the corresponding 19-th degree polynomial through the data points (x, y) were then used to compute and plot the values of this polynomial for 1001 equally-spaced points within the initial range of x . The corresponding result is shown in Figure 1 and the computed coefficients c are given as well.

The code used is given by:

```

127 # 2a
128
129 def q2a(x:np.array, y:np.array, xx:np.array) -> np.ndarray:
130     """Constructing Vandermonde matrix V from data x, LU-decomposing V and solving Vc=y
131     for c.
132     Returns c and the values of the polynomial, yya and ya, corresponding to xx and x.
133     """
134     # Constructing V
135     V = np.zeros((x.shape[0], x.shape[0]), dtype=np.float64)
136     for j in range(V.shape[0]):
137         V[:, j] = np.power(x, j)
138     # LU-decomposing V
139     LU_V = LU_decomposition(V.copy())
140     # Solving for c
141     c0 = solving_system_with_LU(LU_V, y.copy())
142     # Finding the polynomial values y
143     yya = polynomial(c0, xx)
144     ya = polynomial(c0, x)
145     return V, LU_V, c0, yya, ya
146
147 V, LU_V, c0, yya, ya = q2a(x, y, xx)
148
149 # Printing the values of c
150 print("The coefficients of the polynomial found with LU-decomposition are")
151 print(c0)
152
153 # Plot of points with absolute difference shown on a log scale (question 2a)
154 fig=plt.figure()
155 gs=fig.add_gridspec(2,hspace=0,height_ratios=[2.0,1.0])
156 axs=gs.subplots(sharex=True,sharey=False)
157 axs[0].plot(x,y,marker='o',linewidth=0)
158 plt.xlim(-1,101)
159 axs[0].set_ylim(-400,600)
160 axs[0].set_ylabel('$y$')
161 axs[1].set_ylim(1e-18,1e4)
162 axs[1].set_ylabel('$|y-y_i|$')
163 axs[1].set_xlabel('$x$')
164 axs[1].set_yscale('log')
165 line,=axs[0].plot(xx,yya,color='orange')
166 line.set_label('Via LU decomposition')
167 axs[0].legend(frameon=False,loc="lower left")
168 axs[1].plot(x,abs(y-ya),color='orange')
169 plt.savefig('my_vandermonde_sol_2a.png',dpi=600)

```

2_vandermonde.matrix.py

```

1 The coefficients of the polynomial found with LU-decomposition are
2 [ 5.21896930e-01  3.53817596e+03 -2.20652879e+03  6.00128574e+02
3  -9.43549171e+01  9.66286617e+00 -6.87646085e-01  3.53596337e-02
4  -1.34619705e-03  3.84953383e-05 -8.32127939e-07  1.35821060e-08
5  -1.65786631e-10  1.48156916e-12 -9.31152862e-15  3.80148410e-17
6  -8.28186823e-20  2.55562475e-23  1.78743492e-25 -1.28961270e-56]

```

vandermonde.matrix.txt

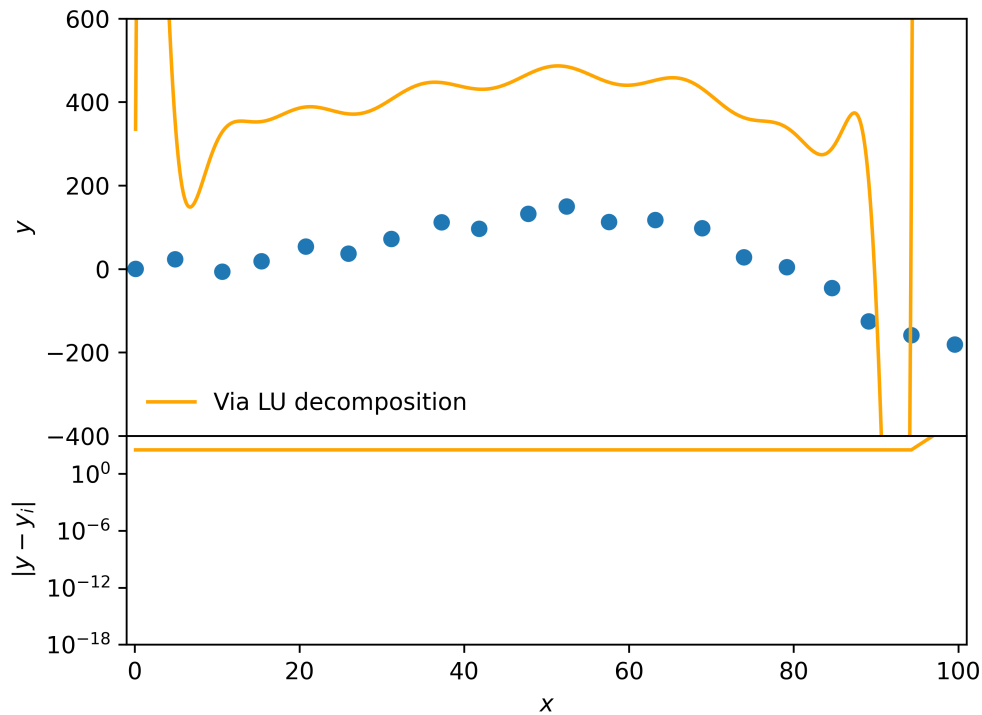


Figure 1: The original data points (x, y) together with the constructed 19-th degree polynomial using the coefficients found by solving the system $Vc = y$ using the LU decomposition of the corresponding Vandermonde matrix are shown. It is clear that the polynomial does not go through the original data points, as it lies way above them. However, the shape of the polynomial does somewhat seem to follow the original data points. It might thus be the case that in particular the offset coefficient c_0 has a large error, causing the polynomial to be shifted this way. The lower panel shows the absolute difference between the polynomial and the data points. As we can see, this offset is about 400 which is very large.

2.2 b

To see whether this polynomial is equal to the Lagrange polynomial, which it should be as they both go through all points and are unique, we compute the Lagrange polynomial using Neville's algorithm on all 20 data points. Although not strictly necessary for this exercise (as $M=20$ and thus all points are used for the interpolation of every point), we have implemented a bisection algorithm to find the M closest points in the data set to each interpolation value x , as this makes for a more generic Neville's algorithm which can be used for other situations and tested for smaller M . The resulting polynomial is shown in Figure 2.

The code used is given by:

```

169 # 2b
170
171 def q2b(x:np.array, y:np.array, xx:np.array, M:np.int64):
172     """Computing the lagrange polynomial through the given data points (x,y) using
173     Neville's algorithm (iff M=x.shape[0]).
174     Returns values yyb and yb of the polynomial corresponding to the points in xx and x
175     respectively."""
176     yyb = neville(M, x, y, xx)[0]
177     yb = neville(M, x, y, x)[0]
178     return yyb, yb
179
180 yyb, yb = q2b(x, y, xx, 20)

```

```

180 # For questions 2b and 2c, add this block
181 line,=axs[0].plot(xx,yyb,linestyle='dashed',color='green')
182 line.set_label('Via Neville\'s algorithm')
183 axs[0].legend(frameon=False,loc="lower left")
184 axs[1].plot(x,abs(y-yb),linestyle='dashed',color='green')
185 plt.savefig('my_vandermonde_sol.2b.png',dpi=600)

```

2_vandermonde_matrix.py

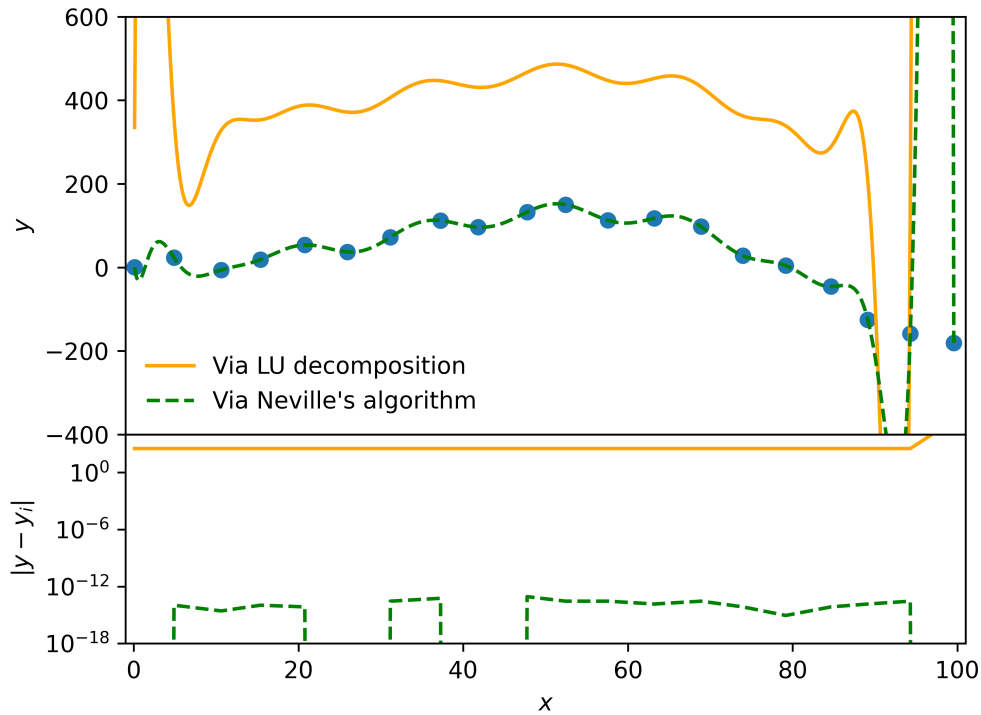


Figure 2: The previous resulting polynomial from LU decomposition is shown together with the Lagrange polynomial found using Neville's algorithm with $M=20$. We see that the Lagrange polynomial does closely follow the original data points, with a small offset shown in the lower panel of about 10^{-14} . The offset even crashes down for some x values, indicating that the Lagrange polynomial very accurately follows the data points. This is to be expected, as Neville's algorithm uses direct linear interpolation on the data points, and when used on the data points themselves, the result is correct up to round off and machine error. The method using the Vandermonde matrix, however, is a more indirect way of computing the polynomial with more intermediate calculations, which does not compute the polynomial values directly but the polynomial coefficients first. As a small deviation in coefficient value can make a large difference in the polynomial values, this explains why the original LU-decomposition method gives polynomial values with a much larger error than the Lagrange polynomial from Neville's algorithm.

2.3 c

We now try to improve the result for LU decomposition by using iterative improvements. From class we know that $V\delta c = \delta y = Vc - y$. Thus we can now solve this system for δc and get $c = c - \delta c$. We use both 1 and 10 iterational improvements on the coefficients and again plot the resulting polynomials. The results are shown in Figure 3.

The code used is given by:

```

187 # 2c

```

```

188
189 def q2c(LU_V:np.ndarray, V:np.ndarray, c:np.array, y:np.array, x:np.array, xx:np.array,
190         num:np.int64) -> np.ndarray:
191     """Computing num iterative improvements on the solution c of the system  $Vc=y$  solved
192     by LU-decomposition.
193     Returns values yyc and yc of the polynomial corresponding to the points in xx and x
194     respectively."""
195     # Iterative improvements
196     for i in range(num):
197         c -= solving_system_with_LU(LU_V, (np.matmul(V,c)-y).copy())
198     # Computing polynomial values
199     yyc = polynomial(c, xx)
200     yc = polynomial(c, x)
201     return yyc, yc
202
203 yyc1, yc1 = q2c(LU_V, V, c0.copy(), y, x, xx, 1)
204 yyc10, yc10 = q2c(LU_V, V, c0.copy(), y, x, xx, 10)
205
206 # For question 2c, add this block too
207 line,=axs[0].plot(xx,yc1,linestyle='dotted',color='red')
208 line.set_label('LU with 1 iteration')
209 axs[1].plot(x,abs(y-yc1),linestyle='dotted',color='red')
210 line,=axs[0].plot(xx,yc10,linestyle='dashdot',color='purple')
211 line.set_label('LU with 10 iterations')
212 axs[1].plot(x,abs(y-yc10),linestyle='dashdot',color='purple')
213 axs[0].legend(frameon=False,loc="lower left")
214 plt.savefig('my_vandermonde_sol_2c.png',dpi=600)

```

2_vandermonde.matrix.py

2.4 d

Lastly, we use timeit to time the execution times of a, b and c (with 10 iterations). We have set timeit's number parameter to 1000, such that we get a more accurate run time estimate without taking more than a minute to execute.

The code used is given by:

```

213 # 2d
214
215 # Timing the previous exercises
216 num = np.int64(1000)
217 print(f"Running 2a {num} times takes", timeit.timeit(lambda: q2a(x,y,xx), number=num), "
218       seconds.")
219 print(f"Running 2b {num} times takes", timeit.timeit(lambda: q2b(x,y,xx,20), number=num)
220       , "seconds.")
221 print(f"Running 2c {num} times takes", timeit.timeit(lambda: q2c(LU_V,V,c0.copy(),y,x,xx
222       ,10), number=num), "seconds.")

```

2_vandermonde.matrix.py

The resulting run times are:

```

1 Running 2a 1000 times takes 1.6017186999961268 seconds.
2 Running 2b 1000 times takes 22.491641599990544 seconds.
3 Running 2c 1000 times takes 1.7273653000011109 seconds.

```

vandermonde.matrix.txt

We see that question b, using Neville's algorithm, uses most computation time by far. This can be explained by the fact that Neville's algorithm computes the value of the polynomial at each of 1001 different x values, directly interpolating between all the x sample points. This requires multiple calculations for each of the 1001 interpolation points. The method using LU decomposition, however, only needs to compute the LU decomposition once and can then compute the 20 coefficients of the polynomial, which can then be used to easily compute the polynomial value for any x. The LU decomposition method is thus more efficient. However, as we have explained before, Neville's algorithm is more accurate in the sense that it has the smallest offset caused by the fact that the method directly interpolates between the

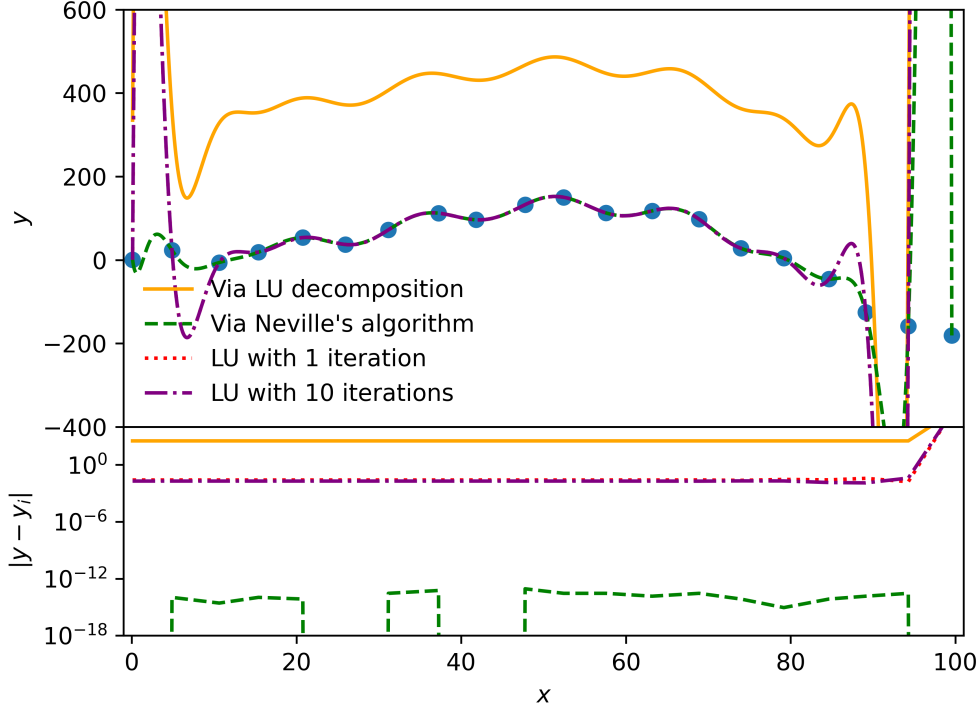


Figure 3: Again, the results from a and b are shown, but the polynomials found using 1 or 10 iterative improvements on the coefficients c are added. We see that the first iterative improvement already greatly improves the polynomial found with LU decomposition: it now follows the data points just as the Lagrange polynomial. However, we do see that the offset is still larger for the improved LU polynomial compared to Neville's Lagrange polynomial. As explained previously, this can be explained by the fact that Neville's algorithm directly computes the polynomial values, while the LU decomposition method computes the coefficients of the polynomial after which those are used to compute the values. This indirect way allows for more error to accumulate. Furthermore, we see that using 10 iterative improvements compared to only 1 does not really improve the result any more. Lastly, we note that all polynomials behave quite normally in the middle of the range of x data points, but more chaotic around the edges of the data points. For example, the LU decomposition polynomials do not even go through the last data point at all. It might be the case that the actual polynomial we are looking for actually behaves more chaotically around the edges, but this can also be a result of the way in which the polynomial is calculated. Especially in the case of Neville's algorithm, the edge points are different from the middle points, because there is no information on what is going on with the polynomial outside the edges.

existing data points. The LU decomposition yield a larger error because it computes the coefficients of the polynomial and not the values of the polynomial directly, causing errors to accumulate and allowing for example the found polynomial to not go through the last data point at all.